



- **1** MapReduce 简介
- **2** MapReduce 原理
- 3 MapReduce 作业管理
- 4 MapReduce 开发案例



MapReduce 简介

▶起源

- 2004年10月Google发表了MapReduce论文
- 设计初衷是解决搜索引擎中大规模网页数据的并行处理
- Hadoop MapReduce是Google MapReduce的开源实现
- MapReduce是Apache Hadoop的核心子项目

> 概念

- 面向离线批处理的分布式计算框架
- 分布式编程模型: MapReduce程序被分为Map(映射)阶段和Reduce(化简)阶段

> 核心思想

- 分而治之,并行计算
- 移动计算, 非移动数据



MapReduce 简介

▶ 特点

- 计算跟着数据走
- •良好的扩展性: 计算能力随着节点数增加, 近似线性增长
- 高容错
- 状态监控
- 适合海量数据的离线批处理
- 降低了分布式编程的门槛



▶ 适用场景

- 数据统计,如网站的PV、UV统计
- 搜索引擎构建索引
- •海量数据查询
- 复杂数据分析算法实现

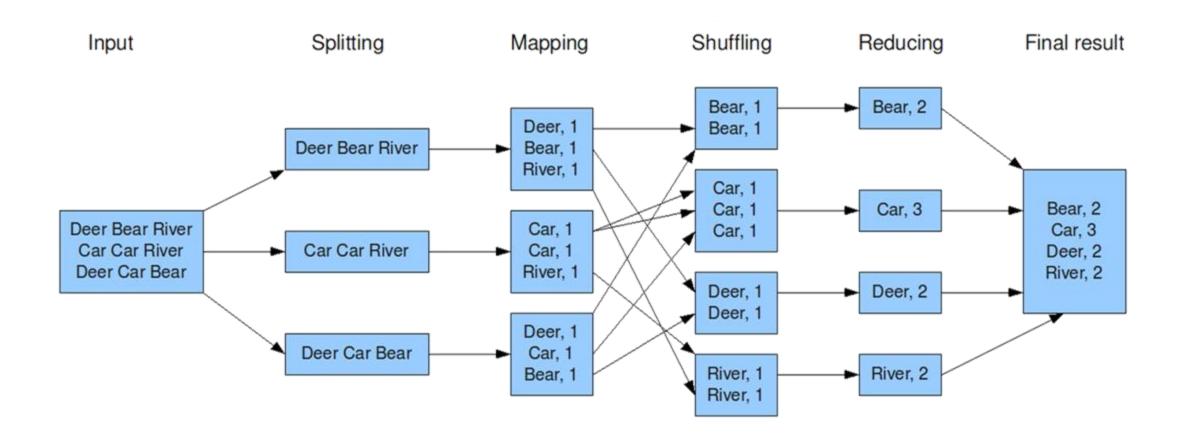
> 不适用场景

- OLAP: 要求毫秒或秒级返回结果
- •流计算:输入数据集是动态的,而MapReduce是静态的
- DAG计算
 - 多个任务之间存在依赖关系,后一个的输入是前一个的输出,构成DAG有向无环图
 - MapReduce很难避免Suffle,造成大量磁盘IO,导致性能较为低下





➤ 示例: WordCount



➤ Job & Task (作业与任务)

- 作业是客户端请求执行的一个工作单元
 - -包括输入数据、MapReduce程序、配置信息
- 任务是将作业分解后得到的细分工作单元
 - 分为Map任务和Reduce任务

➤ Split (切片)

- 输入数据被划分成等长的小数据块,称为输入切片(Input Split),简称切片
- Split是逻辑概念, 仅包含元数据信息, 如数据的起始位置、长度、所在节点等
- 每个Split交给一个Map任务处理,Split的数量决定Map任务的数量
- Split大小
 - 默认等于HDFS Block大小
 - Split的划分方式由程序设定, Split与HDFS Block没有严格的对应关系
 - Split越小,Map任务越多,并发度越高,但开销也越大; Split越大,任务越少,并发度降低TRANSMARP

➤ Map阶段 (映射)

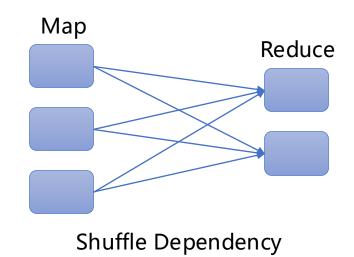
- 由若干Map任务组成,任务数量由Split数量决定
- 输入: Split切片 (key-value)
- 输出:中间计算结果(key-value)

➤ Reduce阶段 (化简)

- 由若干Reduce任务组成,任务数量由程序指定
- 输入: Map阶段输出的中间结果(key-value)
- 输出: 最终结果 (key-value)

➤ Shuffle阶段 (混洗)

- Shuffle是Map和Reduce之间的<mark>强依赖关系(Shuffle依赖</mark>)导致的,即每个Reduce的输入依赖于 所有Map的输出
- Map和Reduce阶段的中间环节(虚拟阶段),分为Map端Shuffle和Reduce端Shuffle



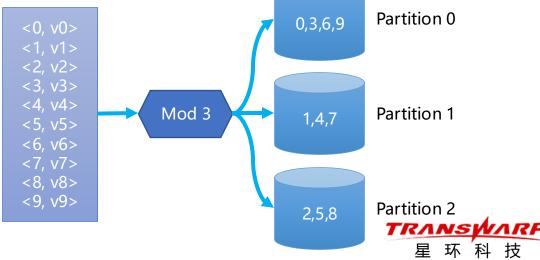


➤ Shuffle阶段(混洗)

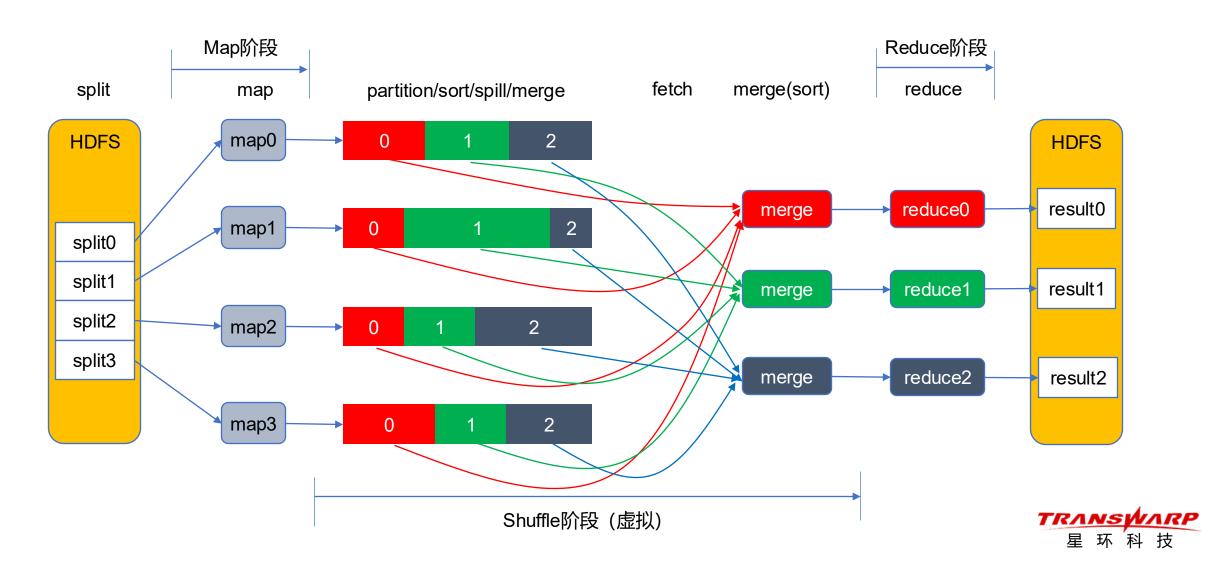
- •包括Partition(分区)、Sort(排序)、Spill(溢写)、Merge(合并)和Fetch(抓取)等工作
- Partition (分区)
 - Reduce任务数量决定了Partition数量, Partition编号 = Reduce任务编号
 - 利用"哈希取模"对Map输出数据分区,即Partition编号 = key hashcode % reduce task num(%为取模)
 - Partition为具有相同编号的Reduce任务供数
 - 哈希取模的作用

✓数据划分:将一个数据集随机分成若干个子集(Hash函数选择不当可能造成数据倾斜)

- ✓数据聚合:将Key相同的数据聚合在一起
- 避免和减少Shuffle是MapReduce程序调优的关键

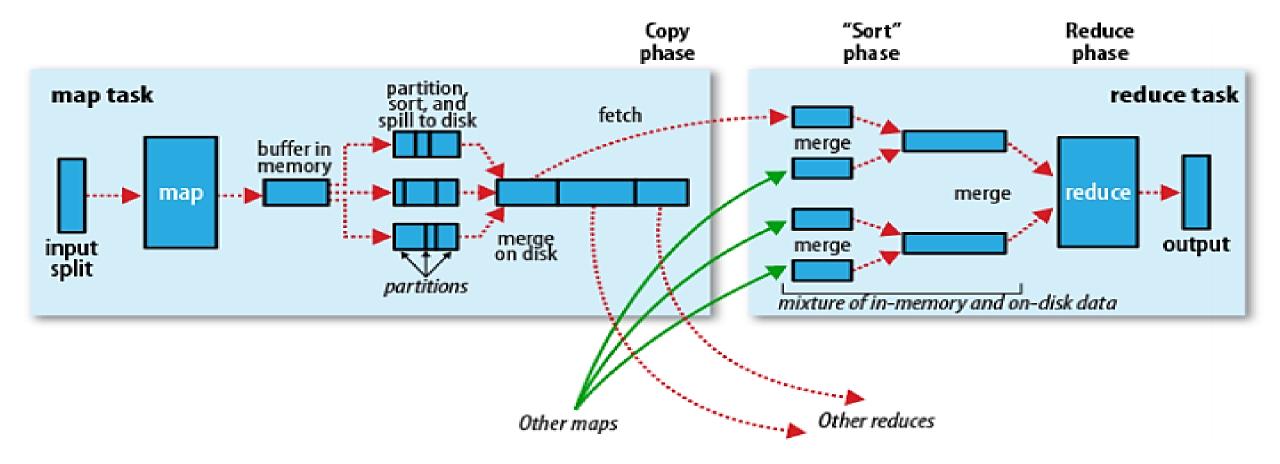


➤ MR执行过程



MR工作机制: Shuffle

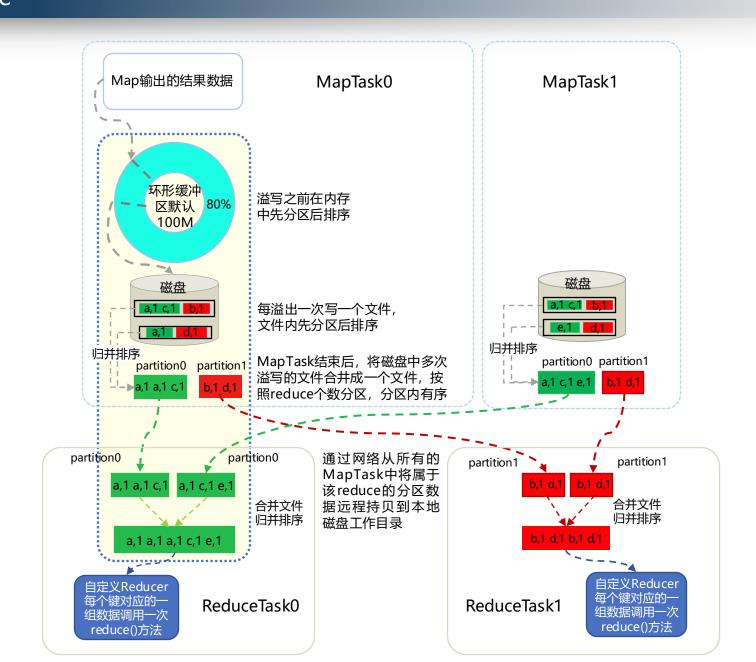
➤ Shuffle详解





MR工作机制: Shuffle

➤ Shuffle详解





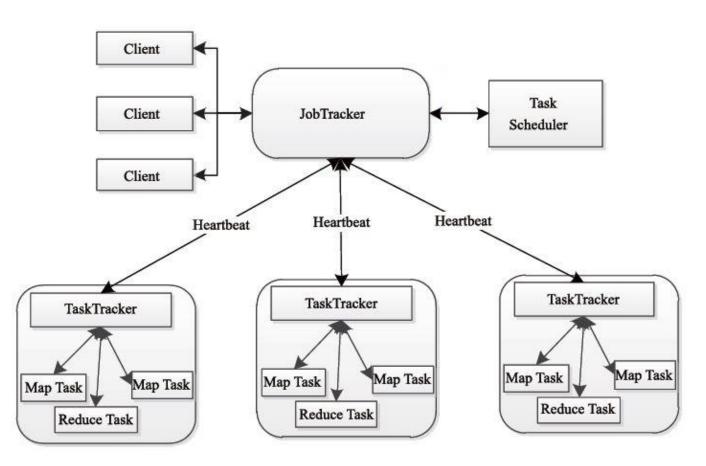
➤ Shuffle详解

- Map端
 - Map任务将中间结果写入环形内存缓冲区Buffer(默认100M),同时进行分区(Partition)和排序(Sort)
 - ✓ 先接"key hashcode % reduce task num"对数据进行分区,分区内再按key排序
 - 当Buffer的数据量达到阈值(默认80%)时,将数据溢写(Spill)到磁盘的一个临时文件中,文件内数据先分区后排序
 - Map任务结束前,将多个临时文件合并(Merge)为一个Map输出文件,文件内数据先分区后排序
- Reduce端
 - Reduce任务从多个Map输出文件中抓取(Fetch)属于自己的分区数据(Partition编号=Reduce任务编号)
 - 对抓取到的分区数据做归并排序,生成一个Reduce输入文件(文件内数据按key排序)
 - ✓ 如果内存缓冲区够大,就直接在内存中完成归并排序,然后落盘
 - ✓ 如果内存缓冲区不够,先将分区数据写到相应的文件中,再通过归并排序合并为一个大文件

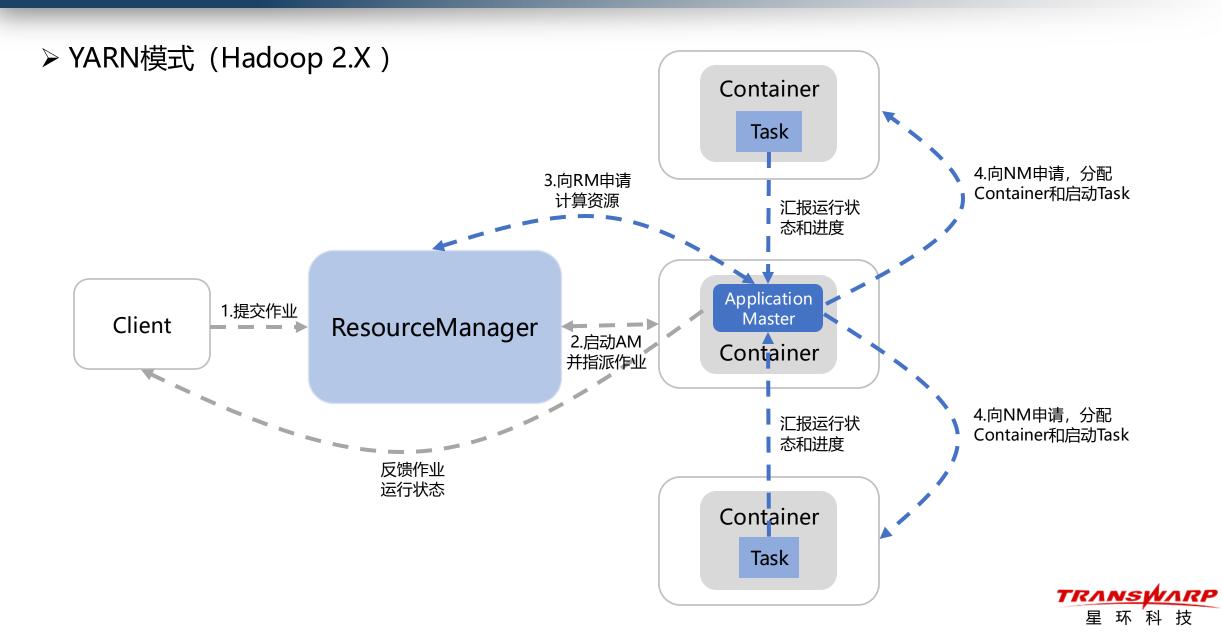


➤ JobTracker/TaskTracker模式 (Hadoop 1.X)

- JobTracker节点 (Master)
 - 调度任务在TaskTracker上运行
 - 若任务失败,指定新TaskTracker重新运行
- TaskTracker节点 (Slave)
 - 执行任务,发送进度报告
- 存在的问题
 - JobTracker存在单点故障
 - JobTracker负载太重(上限4k节点)
 - JobTracker缺少对资源的全面管理
 - TaskTracker对资源的描述过于简单
 - 源码难于理解









1.3 作业管理

▶ 提交作业

```
# hadoop jar {jarFile} [mainClass] args
-jarFile: MapReduce运行程序的jar包
-mainClass: jar包中main函数所在类的类名
-args: 程序调用需要的参数,如输入输出路径
```

> 查看作业

```
# sudo –u yarn application -list
```

> 终止作业

```
# sudo –u yarn application -kill {application_id}
```



▶ 示例: 提交作业

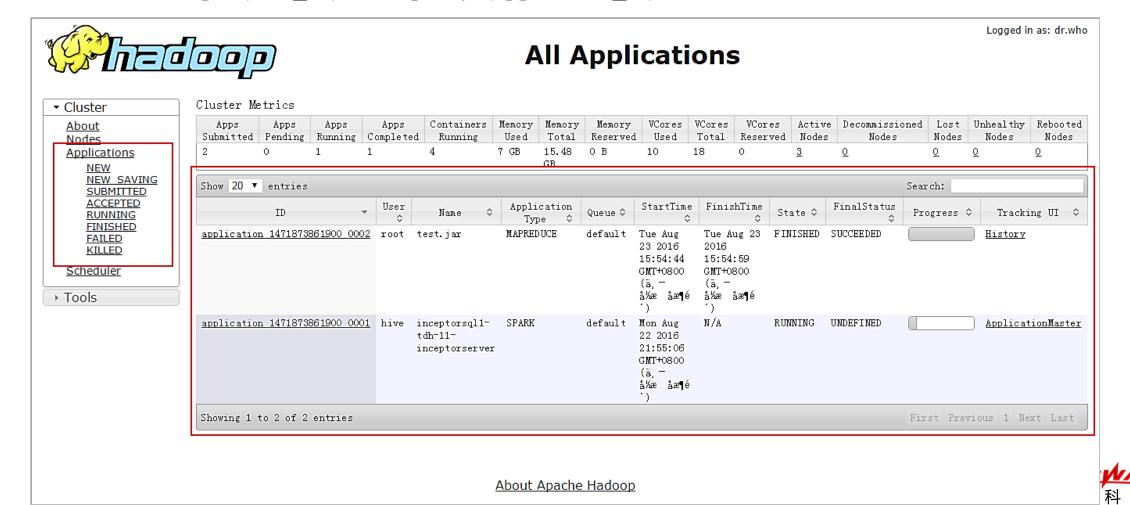
hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-example.jar pi 10 10

```
t3126poc4:~ # hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 10 10
Number of Maps = 10
Samples per Map = 10
Wrote input for Map #0
Wrote input for Map #1
Wrote input for Map #2
Wrote input for Map #3
Wrote input for Map #4
Wrote input for Map #5
Wrote input for Map #6
Wrote input for Map #7
Wrote input for Map #8
Wrote input for Map #9
Starting Job
2016-05-10 14:09:58,250 INFO client.RMProxy: Connecting to ResourceManager at t3126poc5/172.16.2.85:8032
2016-05-10 14:09:58,834 INFO input.FileInputFormat: Total input paths to process : 10
2016-05-10 14:09:58,915 INFO mapreduce.JobSubmitter: number of splits:10
2016-05-10 14:09:59,188 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1462786145119 0002
2016-05-10 14:09:59,453 INFO impl.YarnClientImpl: Submitted application application 1462786145119 0002
2016-05-10 14:09:59,498 INFO mapreduce.Job: The url to track the job: http://t3126poc5:8088/proxy/application 1462786145119 0002/
2016-05-10 14:09:59,499 INFO mapreduce.Job: Running job: job 1462786145119 0002
2016-05-10 14:10:05,641 INFO mapreduce.Job: Job job 1462786145119 0002 running in uber mode : false
2016-05-10 14:10:05.644 INFO mapreduce.Job: map 0% reduce 0%
```



> 作业监控

• Web监控: http://{AM_IP}:8088/proxy/{application_id}/



1.3 作业管理

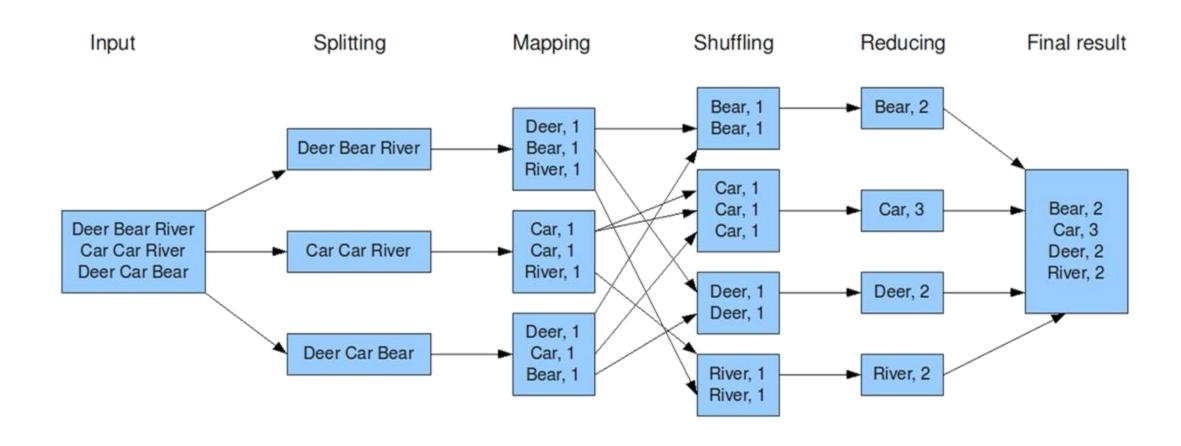
- > 作业诊断
 - 配置参数: yarn.nodemanager.log-dirs
 - MapReduce运行日志目录
 - 默认值为/mnt/disk*/hadoop/yarn/
 - 根据运行出错信息,到指定节点下分析日志

t3126poc5:~ # ls /mnt/disk2/hadoop/yarn/logs/application_1462783245088_0002/container_1462783245088_0002_01_000002/ stderr stdout syslog





> WordCount:



▶ 代码:

WordCountMapper

```
public class WordCountMapper extends Mapper < LongWritable, Text, Text, IntWritable > {
    @Override
    protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        String[] words = StringUtils.split(value.toString(), '');

    for (String word : words) {
        context.write(new Text(word), new IntWritable(1));
    }
}
```

≻ 代码:

WordCountReducer

```
public class WordCountReducer extends Reducer < Text, IntWritable, Text, IntWritable> {
    @Override
    protected void reduce(Text key, Iterable < IntWritable > values, Context context) throws IOException, InterruptedException {
    int sum = 0;
    for (IntWritable i : values) {
        sum = sum + i.get();
    }
    context.write(key, new IntWritable(sum));
}
```

➤ 代码: WordCountJob

```
Configuration conf = new Configuration();
Job job = Job. getInstance(conf);
job.setJobName("WordCount");
job.setJarByClass(WordCountJob.class);
// 指定Map、Reduce实现类
job.setMapperClass(WordCountMapper.class);
job.setReducerClass(WordCountReducer.class);
// 指定Mapper输出时 Key Value数据类型
job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);
#词频分析 测试数据 在HDFS中的存放位置
Path inputPath = new Path("/tmp/dir4test/wordcount/input");
FileInputFormat.addInputPath(job, inputPath);
// 词频分析 分析结果 HDFS输出位置
Path outputPath = new Path("/tmp/dir4test/wordcount/output");
FileOutputFormat. setOutputPath(job, outputPath);
boolean flag = job.waitForCompletion(true);
System. out.println(flag? "Success!": "Error~");
```

➤ 代码: WordCountJob

```
Configuration conf = new Configuration();
Job job = Job. getInstance(conf);
job.setJobName("WordCount");
job.setJarByClass(WordCountJob.class);
// 指定Map、Reduce实现类
job.setMapperClass(WordCountMapper.class);
job.setReducerClass(WordCountReducer.class); // Combiner
                                                   job.setCombinerClass(WordCountMapper.class);
// 指定Mapper输出时 Key Value数据类型
job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);
// 词频分析 测试数据 在HDFS中的存放位置
Path inputPath = new Path("/tmp/dir4test/wordcount/input");
FileInputFormat.addInputPath(job, inputPath);
// 词频分析 分析结果 HDFS输出位置
Path outputPath = new Path("/tmp/dir4test/wordcount/output");
FileOutputFormat. setOutputPath(job, outputPath);
boolean flag = job.waitForCompletion(true);
System. out.println(flag? "Success!": "Error~");
```

▶ 二次排序:

• 需求: 现有海量的历史天气数据, 统计出不同月份中温度最高的日期。

• 数据:

```
1949-10-01 14:21:02
                            34c
1949-10-02 14:01:02
                            36c
1950-01-01 11:21:02
                            32c
1950-10-01 12:21:02
                            37c
1951-12-01 12:21:02
                            23c
1950-10-02 12:21:02
                            41c
1950-10-03 12:21:02
                            27c
1951-07-01 12:21:02
                            45c
1951-07-02 12:21:02
                            46c
1951-07-03 12:21:03
                            47c
... ...
```



▶代码: 自定义组合Key

```
public class Weather implements WritableComparable
                                                     <Weather> {
  private int year;
  private int month;
  private int day;
  private int wd; // 温度
  public int getYear() {
    return year; }
  public void setYear(int year) {
    this.year = year; }
  //部分set、get代码略...
  @Override
  public void write(DataOutput out) throws IOException {
    out.writeInt(year);
    out.writeInt(month);
    out.writeInt(day);
    out.writeInt(wd);
```

```
@Override
 public void readFields(DataInput in) throws IOException {
    this.year = in.readInt();
    this.month = in.readInt();
    this.day = in.readInt();
    this.wd = in.readInt();
 @Override
 public int compareTo(Weather w) {
    int c1 = Integer.compare(this.year, w.getYear());
   if(c1 == 0)
      int c2 = Integer.compare(this.month, w.getMonth());
      if(c2 == 0)
         return Integer.compare(this.wd, w.getWd());
      return c2;
    return c1;
```

▶代码: 自定义Sort

```
public class TQSort extends WritableComparator {
  public TQSort() {
    super(Weather.class, true);
  @Override
  public int compare(WritableComparable a, WritableComparable b) {
    Weather w1 = (Weather)a;
    Weather w2 = (Weather)b;
    int c1 = Integer.compare(w1.getYear(), w2.getYear());
    if(c1 == 0) {
       int c2 = Integer.compare(w1.getMonth(), w2.getMonth());
       if(c2 == 0)
         return -Integer.compare(w1.getWd(), w2.getWd());
       return c2;
    return c1;
```

▶代码: 自定义Group

```
public class TQGroup extends WritableComparator {
  public TQGroup() {
    super(Weather.class, true);
  @Override
  public int compare(WritableComparable a, WritableComparable b) {
    Weather w1 = (Weather)a;
    Weather w2 = (Weather)b;
    int c1 = Integer.compare(w1.getYear(), w2.getYear());
    if(c1 == 0) {
       int c2 = Integer.compare(w1.getMonth(), w2.getMonth());
       return c2;
    return c1;
```

▶代码: 自定义Partition

```
public class TQPartition extends HashPartitioner < Weather, IntWritable > {
    @Override
    public int getPartition(Weather key, IntWritable value, int numReduceTasks) {
        return (key.getYear() - 1949 ) % numReduceTasks;
        // return super.getPartition(key, value, numReduceTasks);
    }
}
```

➤ 代码: Mapper

```
public class TQMapper extends Mapper LongWritable, Text, Weather, IntWritable \{
  @Override
  protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
    String[] strs = StringUtils.split(value.toString(), '\t');
    SimpleDateFormat sdf = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss");
    Calendar cal = Calendar.getInstance();
    try {
       cal.setTime(sdf.parse(strs[0]));
       Weather w = new Weather():
       w.setYear(cal.get(Calendar.YEAR));
       w.setMonth(cal.get(Calendar.MONTH) + 1);
       w.setDay(cal.get(Calendar.DAY OF MONTH));
       int wd = Integer.parseInt(strs[1].substring(0, strs[1].lastIndexOf("c")));
       w.setWd(wd);
       context.write(w, new IntWritable(wd));
    } catch (ParseException e) {
```

➤ 代码: Reducer

```
public class TQReducer extends Reducer<Weather, IntWritable, Text, NullWritable> {
  @Override
  protected void reduce(Weather weather, Iterable<IntWritable> iterable, Context context)
       throws IOException, InterruptedException {
     int flag = 0;
     for (IntWritable i : iterable) {
       flag++;
       if(flag > 2) {
          break;
       String msg = weather.getYear() + "-" + weather.getMonth() + "-" + weather.getDay() +
            <u>"-"</u> + i.get();
       context.write(new Text(msg), NullWritable.get());
```

➤ 代码: Job

```
Configuration conf = new Configuration();
Job job = Job. getInstance(conf);
job.setJobName("TQ sort");
job.setJarByClass(TQJob.class);
// 指定Map、Reduce实现类
job.setMapperClass(TQMapper.class);
job.setReducerClass(TQReducer.class);
// 指定Mapper输出时 Key Value数据类型
job.setMapOutputKeyClass(Weather.class);
job.setMapOutputValueClass(IntWritable.class);
// 指定自定义Partition、Sort、Group
job.setPartitionerClass(TQPartition.class);
job.setSortComparatorClass(TQSort.class);
job.setGroupingComparatorClass(TQGroup.class);
// 指定Reduce Task 数量
job.setNumReduceTasks(3);
```

```
// 二次排序 测试数据 在HDFS中的存放位置
Path inputPath = new Path("/tmp/dir4test/weather/input");
FileInputFormat.addInputPath(job, inputPath);
// 二次排序 分析结果 HDFS输出位置
Path outputPath = new Path("/tmp/dir4test/weather/output");
FileOutputFormat.setOutputPath(job, outputPath);
boolean flag = job.waitForCompletion(true);
System. out.println(flag? "Success!": "Error~");
```



温故知新

- 简述MR Split与HDFS Block的关系。
- 为什么MapReduce要求输入输出必须是key-value键值对?
- · 简述Shuffle的工作原理。
- · 从编程模型的视角, MapReduce有哪些优缺点?
- · 简述"哈希取模"在MapReduce中的作用。

